

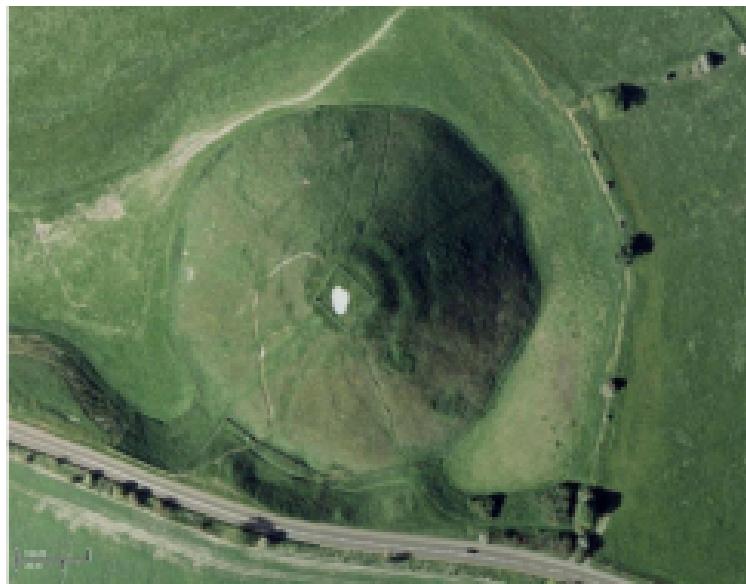
RECAP

- Where do we look in the image, to describe an image locally?..



RECAP

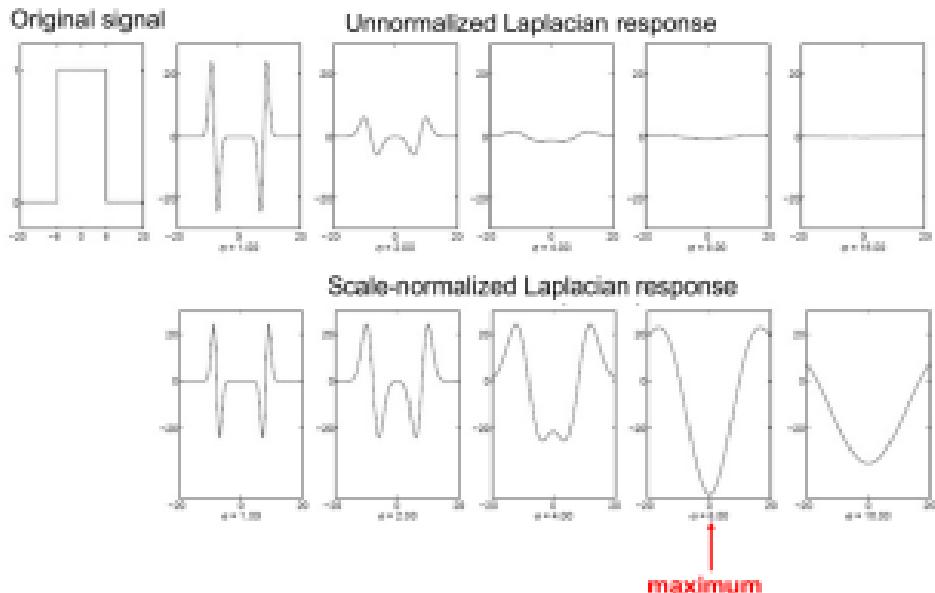
- Where do we look in the image, to describe an image locally?..



Informative regions: along edges, around corners/keypoints, blobs.

- How do we detect these informative regions in the image?..

■ How do we detect these informative regions in the image?..



RECAP

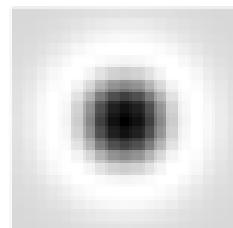
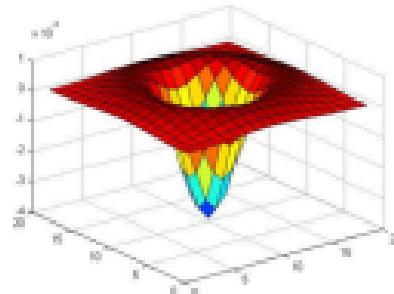
- What is the Laplacian? And in what way is it normalized?

RECAP

- What is the Laplacian? And in what way is it normalized?

Blob detection in 2D

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



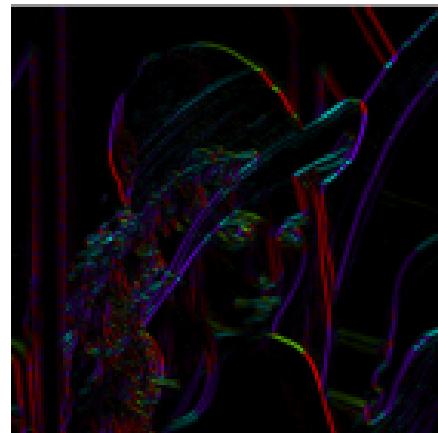
Scale-normalized: $\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$

RECAP

- What image information do we extract to describe an image (locally)?...
- Image edge orientations.



Lena



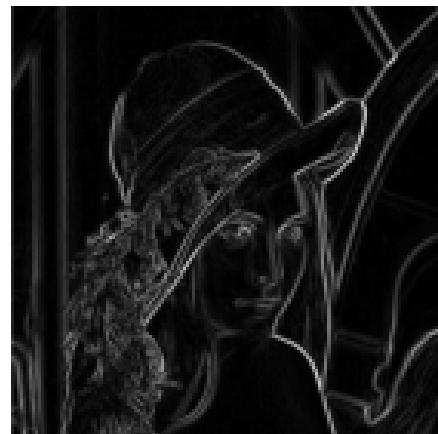
Gradient orientation

RECAP

- What image information do we extract to describe an image (locally)?...
- Image edge magnitude.

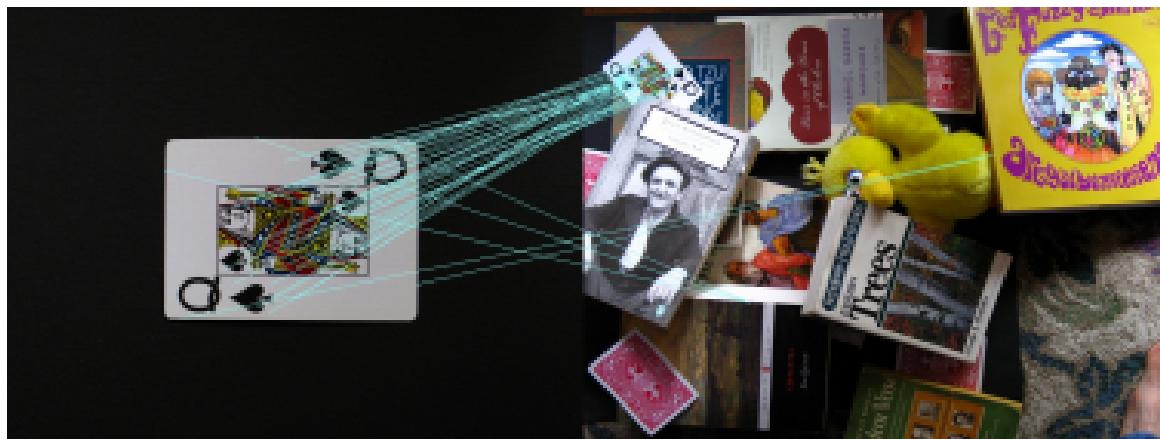


Lena

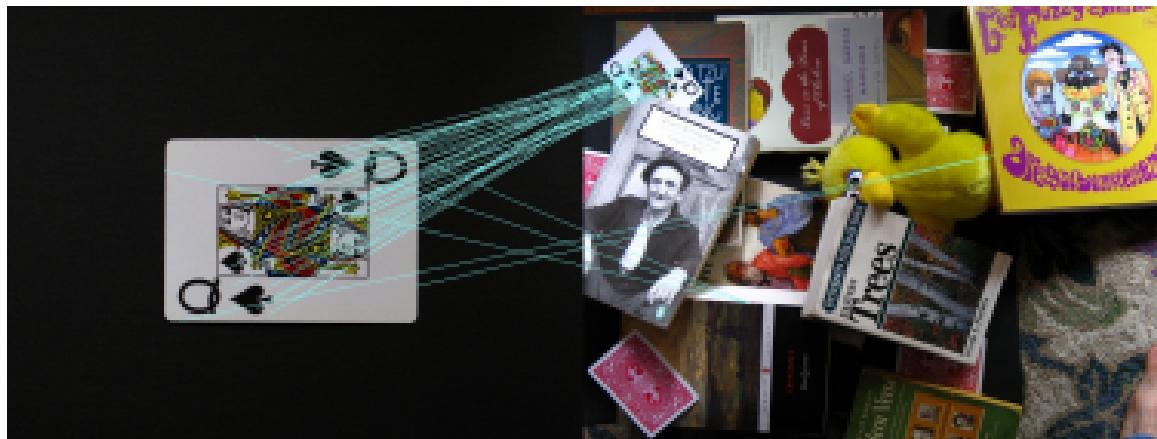


Gradient magnitude

SIFT: MOTIVATION



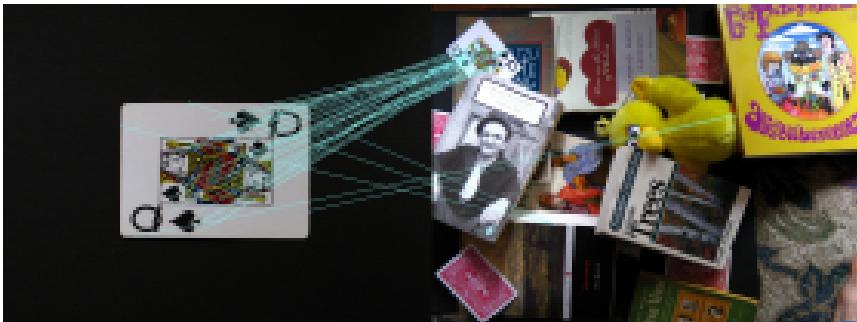
SIFT: MOTIVATION



Find local features — object dependent, focused on the consistent part, ignore background.

SIFT: MOTIVATION

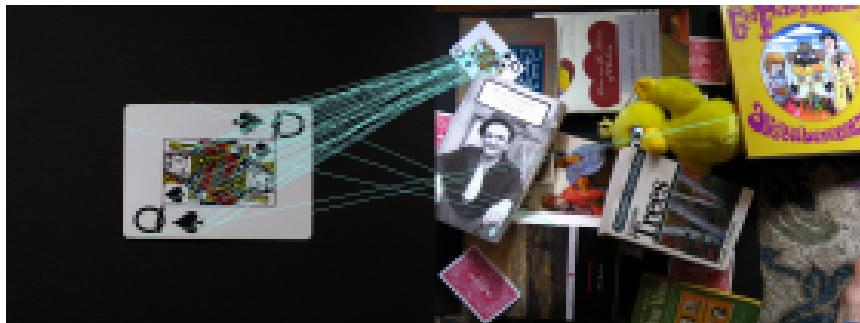
- Image matching...



- Desirable characteristics...

SIFT: MOTIVATION

■ Image matching...



■ Desirable characteristics...

- ▶ Scale invariance.
- ▶ Rotation invariance.
- ▶ Illumination invariance.
- ▶ Viewpoint invariance.

SIFT: METHOD OVERVIEW

1. CONSTRUCTING A SCALE-SPACE.
2. LAPLACIAN OF GAUSSIAN APPROXIMATION.
3. FINDING KEYPOINTS.
4. REMOVE UNINFORMATIVE KEYPOINTS.
5. ASSIGN ORIENTATION TO KEYPOINTS.
6. GENERATE A SIFT DESCRIPTOR.

CONSTRUCTING A SCALE-SPACE

Why do we need a scale-space?...



CONSTRUCTING A SCALE-SPACE

How do we construct a scale-space?...



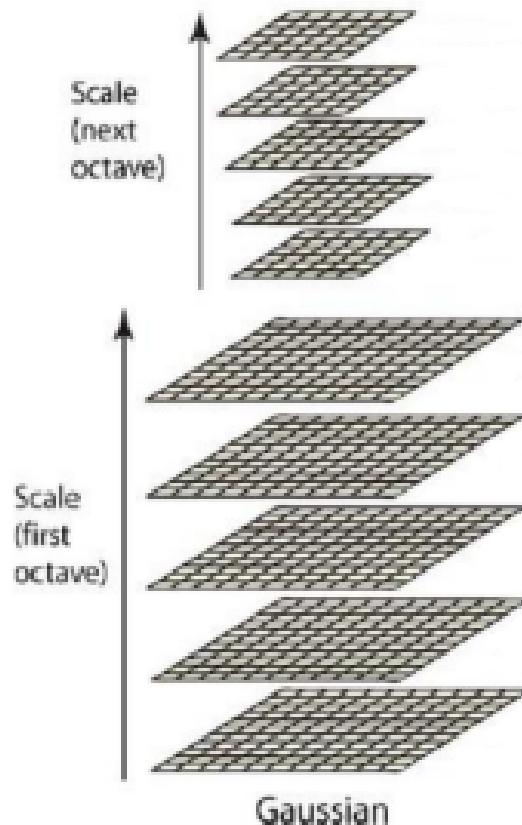
CONSTRUCTING A SCALE-SPACE

- Gaussian kernel is the only that is mathematically proven to construct a scale space without spurious resolution [Lindberg, 1994]

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y), \quad (1)$$

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2} \quad (2)$$

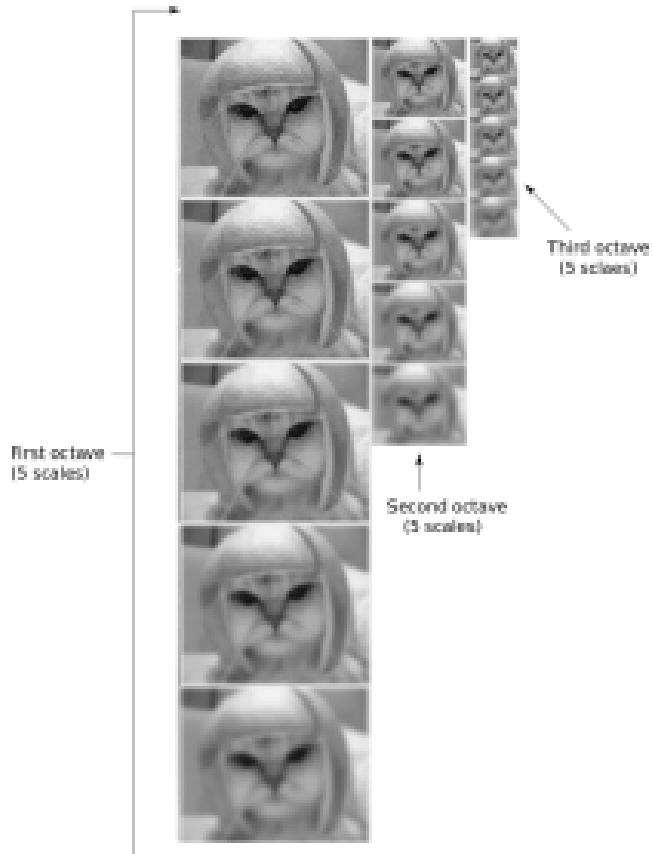
SIFT SCALE-SPACE: OCTAVES



SIFT SCALE-SPACE: OCTAVES

1. Take the original image and progressively blur it $5\times$ with increasing σ .
2. Resize the image to its half and repeat the process.
3. Each "octave" has 5 image.
4. Each image in an octave has a different "scale" (amount of blur).

SIFT SCALE-SPACE: OCTAVES



SIFT: METHOD OVERVIEW

1. CONSTRUCTING A SCALE-SPACE.
2. LAPLACIAN OF GAUSSIAN APPROXIMATION.
3. FINDING KEYPOINTS.
4. REMOVE UNINFORMATIVE KEYPOINTS.
5. ASSIGN ORIENTATION TO KEYPOINTS.
6. GENERATE A SIFT DESCRIPTOR.

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of a 2 D function $f(x, y)$...

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of a 2 D function $f(x, y)$...

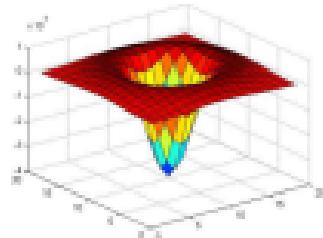
$$\nabla^2 f(x, y) = \frac{\partial^2}{\partial x^2} f(x, y) + \frac{\partial^2}{\partial y^2} f(x, y) \quad (3)$$

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of a 2 D function $f(x, y)$...

$$\nabla^2 f(x, y) = \frac{\partial^2}{\partial x^2} f(x, y) + \frac{\partial^2}{\partial y^2} f(x, y) \quad (3)$$

- Laplacian of Gaussian:



$$\nabla^2 G(x, y, \sigma) = \left(\frac{x^2 + y^2 - \sigma^2}{\sigma^4} \right) e^{-(x^2 + y^2)/2\sigma^2}$$

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of Gaussian computation on images...

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of Gaussian computation on images...

$$\nabla^2(f(x, y) \otimes G(x, y)) = \underbrace{\nabla^2 G(x, y)}_{\text{Laplacian of Gaussian-filtered image}} \otimes f(x, y)$$

Laplacian of
Gaussian-filtered image

Laplacian of Gaussian (LoG)
-filtered image

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of Gaussian computation on images...

$$\nabla^2(f(x, y) \otimes G(x, y)) = \underbrace{\nabla^2 G(x, y)}_{\text{Laplacian of Gaussian-filtered image}} \otimes f(x, y)$$

Laplacian of
Gaussian-filtered image

Laplacian of Gaussian (LoG)
-filtered image

- It is useful for...

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of Gaussian computation on images...

$$\nabla^2(f(x, y) \otimes G(x, y)) = \nabla^2G(x, y) \otimes f(x, y)$$

Laplacian of
Gaussian-filtered image

Laplacian of Gaussian (LoG)
-filtered image

- It is useful for...

- ▶ Locating edges and corners in an image.
- ▶ Edges and corners are good for finding keypoints.

LAPLACIAN OF GAUSSIAN APPROXIMATION

- Laplacian of Gaussian computation on images...

$$\nabla^2(f(x, y) \otimes G(x, y)) = \nabla^2G(x, y) \otimes f(x, y)$$

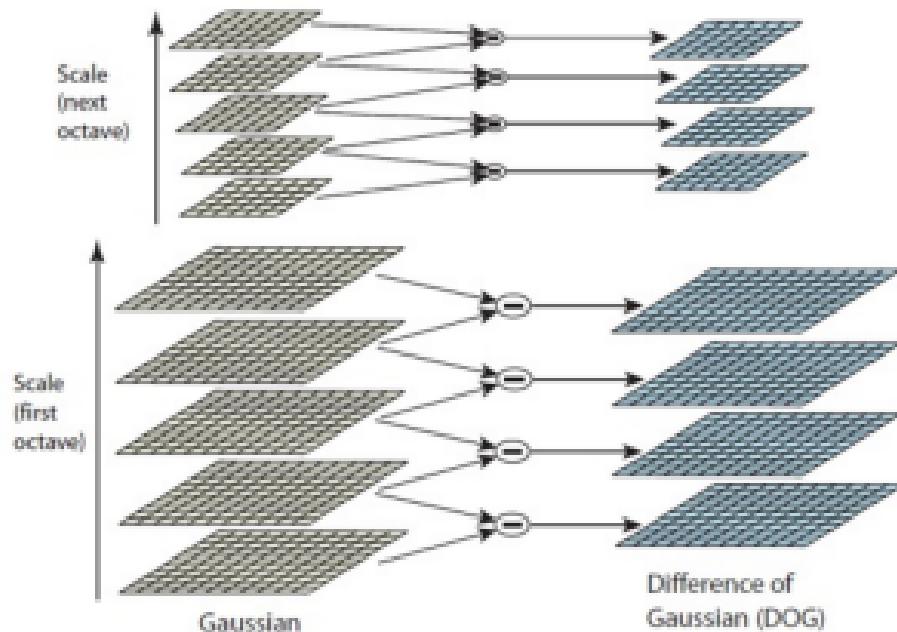
Laplacian of
Gaussian-filtered image

Laplacian of Gaussian (LoG)
-filtered image

- It is useful for...
 - ▶ Locating edges and corners in an image.
 - ▶ Edges and corners are good for finding keypoints.
- But second order derivatives are expensive to compute.. and we already have the blurred images.

LAPLACIAN OF GAUSSIAN APPROXIMATION

We can approximate the Laplacian of Gaussian (LoG) with a difference of Gaussian (DoG). Why?



LAPLACIAN OF GAUSSIAN APPROXIMATION

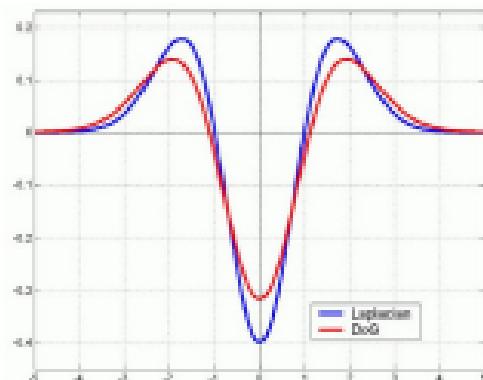
LoG can be approximate by a Difference of two Gaussians (DoG) at different scales.

$$\text{LoG} = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



To approximate a LoG for some σ , we need two gaussians with sigmas: $\sqrt{\sigma^2 + \Delta t}$ and $\sqrt{\sigma^2 - \Delta t}$.

LAPLACIAN OF GAUSSIAN APPROXIMATION

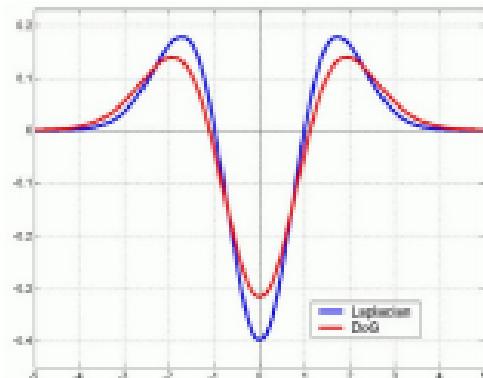
LoG can be approximate by a Difference of two Gaussians (DoG) at different scales.

$$\text{LoG} = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



To approximate a LoG for some σ , we need two gaussians with sigmas: $\sqrt{\sigma^2 + \Delta t}$ and $\sqrt{\sigma^2 - \Delta t}$.

LAPLACIAN OF GAUSSIAN APPROXIMATION: BENEFITS

- Images blurred with a Gaussian kernel are not scale independent. Why?...

LAPLACIAN OF GAUSSIAN APPROXIMATION: BENEFITS

- Images blurred with a Gaussian kernel are not scale independent. Why?...

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2} \quad (4)$$

LAPLACIAN OF GAUSSIAN APPROXIMATION: BENEFITS

- Images blurred with a Gaussian kernel are not scale independent. Why?...

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2} \quad (4)$$

- The "scale-invariant" Laplacian should look like this:
 $\sigma^2 \nabla^2 G(x, y, \sigma)$.

LAPLACIAN OF GAUSSIAN APPROXIMATION: BENEFITS

- Images blurred with a Gaussian kernel are not scale independent. Why?...

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2} \quad (4)$$

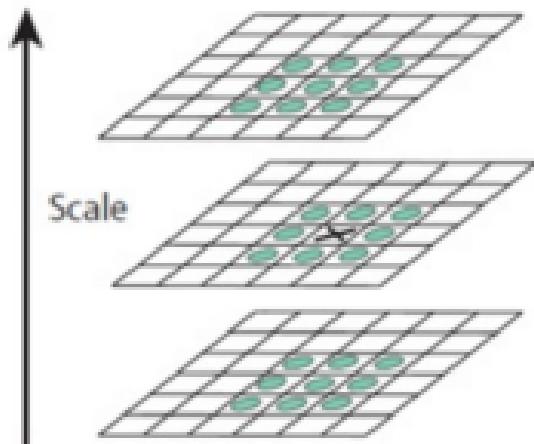
- The "scale-invariant" Laplacian should look like this:
 $\sigma^2 \nabla^2 G(x, y, \sigma)$.
- And that is the result of DoG.

SIFT: METHOD OVERVIEW

1. CONSTRUCTING A SCALE-SPACE.
2. LAPLACIAN OF GAUSSIAN APPROXIMATION.
3. FINDING KEYPOINTS.
4. REMOVE UNINFORMATIVE KEYPOINTS.
5. ASSIGN ORIENTATION TO KEYPOINTS.
6. GENERATE A SIFT DESCRIPTOR.

FINDING KEYPOINTS

Find the local minima/maxima by looking at all neighbors of all pixels (across scales also):



To find keypoints at subpixel location, an image Taylor expansion around the minima/maxima points is used.

FINDING KEYPOINTS

- How many extrema (minima/maxima) images we have in the end per octave?...

FINDING KEYPOINTS

- How many extrema (minima/maxima) images we have in the end per octave?...
- 5 scales \rightarrow 4 DoG images \rightarrow 2 extrema images (the first and last DoG layer do not have sufficient neighbors, thus no extrema there).

